

# AIML for Finance - Using Quantmod to Import Financial Data

*Brian Clark*

*Fall 2020*

## Contents

<b>1</b>	<b>Overview</b>	<b>2</b>
<b>2</b>	<b>Downloading Stock Data into R using “quantmod”</b>	<b>2</b>
<b>3</b>	<b>Macro Data</b>	<b>6</b>
<b>4</b>	<b>Time-Series Data in R</b>	<b>7</b>
<b>5</b>	<b>References</b>	<b>14</b>

---

# 1 Overview

This document is broken down into three parts. The first describes how to use R’s “quantmod” package to download stock data. This part is based on a vignette put together and maintained by Majeed Simaan who is a Lally School (RPI) PhD and is now an Assistant Professor of Finance at Stevens Institute of Technology. A link to his github page with the details of his vignette is [here](#). The second part includes examples of how to bring in other relevant financial data. The third part shows some useful time-series commands.

## 2 Downloading Stock Data into R using “quantmod”

This section is based on Majeed Simaan’s R vignette. Please refer to his original file as his vignette also contains instructions and examples of how to download and manage options data along with macroeconomic data hosted by the **Federal Reserve Bank of St. Louis**.

Prior to starting, you need to install the “quantmod” package. For this example, we download daily equity data for six stocks starting on January 1, 1980.

```
rm(list=ls())

# install.packages("quantmod")
library("quantmod")

# Define some companies:
tics <- c("AAPL", "GE", "TWTR", "BAC", "RAD", "PFE")
P.list <- lapply(tics, function(tic)
  get(getSymbols(tic, from = "1980-01-01")))

sapply(P.list, nrow)
```

```
## [1] 9862 10102 1562 10102 10050 10102
```

A few comments are in order for the above block of code. The tics object obviously contains the ticker symbols you want to use. They correspond to the tickers on Yahoo Finance so if you want other companies you can search for the tickers [here](#). We will stick to Majeed’s example for now.

The “lapply” function is another way to write a for loop in R. For example, P.list could also be defined (equivalently) as follows:

```
# The following block is equivalent to the lapply command::
P.list2 <- list(NA)
n <- length(tics)
for (i in 1:n){
  P.list2[[i]] <- get(getSymbols(tics[i], from = "1980-01-01"))
}

# ...is equivalent to...

# P.list <- lapply(tics, function(tic)
#   get(getSymbols(tic, from = "1980-01-01")))

# Printing the number of rows could also be accomplished in a loop:
for (i in 1:n){
  print(nrow(P.list2[[i]]))
}
```

```
## [1] 9862
## [1] 10102
```

```
## [1] 1562
## [1] 10102
## [1] 10050
## [1] 10102
```

```
sapply(P.list,nrow)
```

```
## [1] 9862 10102 1562 10102 10050 10102
```

Note that each of the stocks has a different number of rows. This is because we are using real data and as such there will be missing values (for example, Twitter was not around in 1980!). Let's print the first and last few observations for each stock:

```
for (i in 1:n){
  print(tics[i])
  print(head(P.list[[i]],4))
  print("-----")
}
```

```
## [1] "AAPL"
##           AAPL.Open AAPL.High AAPL.Low AAPL.Close AAPL.Volume
## 1980-12-12  0.513393  0.515625 0.513393  0.513393   117258400
## 1980-12-15  0.488839  0.488839 0.486607  0.486607   43971200
## 1980-12-16  0.453125  0.453125 0.450893  0.450893   26432000
## 1980-12-17  0.462054  0.464286 0.462054  0.462054   21610400
##           AAPL.Adjusted
## 1980-12-12      0.407747
## 1980-12-15      0.386473
## 1980-12-16      0.358108
## 1980-12-17      0.366972
## [1] "-----"
## [1] "GE"
##           GE.Open  GE.High  GE.Low GE.Close GE.Volume GE.Adjusted
## 1980-01-02 1.014123 1.016627 0.976563 0.976563  7433000  0.004071
## 1980-01-03 0.976563 0.991587 0.959034 0.989083  9185200  0.004124
## 1980-01-04 0.999099 1.024139 0.999099 1.021635  8556200  0.004259
## 1980-01-07 1.021635 1.071715 1.014123 1.056691 10518100  0.004405
## [1] "-----"
## [1] "TWTR"
##           TWTR.Open TWTR.High TWTR.Low TWTR.Close TWTR.Volume
## 2013-11-07      45.10      50.09      44.00      44.90   117701600
## 2013-11-08      45.93      46.94      40.69      41.65   27925300
## 2013-11-11      40.50      43.00      39.40      42.90   16113900
## 2013-11-12      43.66      43.78      41.83      41.90    6316700
##           TWTR.Adjusted
## 2013-11-07      44.90
## 2013-11-08      41.65
## 2013-11-11      42.90
## 2013-11-12      41.90
## [1] "-----"
## [1] "BAC"
##           BAC.Open BAC.High  BAC.Low BAC.Close BAC.Volume BAC.Adjusted
## 1980-01-02      0 1.718750 1.687500  1.687500    36000  0.130777
## 1980-01-03      0 1.687500 1.671875  1.671875    31200  0.129566
## 1980-01-04      0 1.671875 1.640625  1.671875    35200  0.129566
## 1980-01-07      0 1.703125 1.656250  1.687500   260000  0.130777
```

```
## [1] "-----"
## [1] "RAD"
##      RAD.Open RAD.High RAD.Low RAD.Close RAD.Volume RAD.Adjusted
## 1980-03-17      0 28.28125 27.03125  27.1875      11400      19.13496
## 1980-03-18      0 26.87500 26.25000  26.5625      27400      18.69507
## 1980-03-19      0 26.87500 25.93750  25.9375      12000      18.25519
## 1980-03-20      0 26.25000 25.62500  25.6250       2900      18.03524
## [1] "-----"
## [1] "PFE"
##      PFE.Open PFE.High PFE.Low PFE.Close PFE.Volume PFE.Adjusted
## 1980-01-02 0.809896 0.809896 0.781250 0.781250    3216000      0.000681
## 1980-01-03 0.781250 0.789063 0.770833 0.781250    2846400      0.000681
## 1980-01-04 0.789063 0.809896 0.789063 0.809896    3316800      0.000706
## 1980-01-07 0.809896 0.820313 0.799479 0.809896    2184000      0.000706
## [1] "-----"
```

```
for (i in 1:n){
  print(tics[i])
  print(tail(P.list[[i]],4))
  print("-----")
}
```

```
## [1] "AAPL"
##      AAPL.Open AAPL.High AAPL.Low AAPL.Close AAPL.Volume
## 2020-01-17    316.27    318.74    315.00    318.73    34454100
## 2020-01-21    317.19    319.02    316.00    316.57    27710800
## 2020-01-22    318.58    319.99    317.31    317.70    25458100
## 2020-01-23    317.92    319.56    315.65    319.23    26071600
##      AAPL.Adjusted
## 2020-01-17        318.73
## 2020-01-21        316.57
## 2020-01-22        317.70
## 2020-01-23        319.23
## [1] "-----"
## [1] "GE"
##      GE.Open GE.High GE.Low GE.Close GE.Volume GE.Adjusted
## 2020-01-17    11.85    11.93    11.76    11.81    46580700      11.81
## 2020-01-21    11.74    11.80    11.61    11.66    43738200      11.66
## 2020-01-22    11.71    11.71    11.35    11.37    61947800      11.37
## 2020-01-23    11.73    11.85    11.60    11.77    73125200      11.77
## [1] "-----"
## [1] "TWTR"
##      TWTR.Open TWTR.High TWTR.Low TWTR.Close TWTR.Volume
## 2020-01-17    33.82    34.39    33.58    34.22    19303500
## 2020-01-21    34.08    34.39    33.87    34.09    17055600
## 2020-01-22    34.29    34.54    33.96    34.02    12851800
## 2020-01-23    33.99    34.36    33.72    33.89    14879100
##      TWTR.Adjusted
## 2020-01-17        34.22
## 2020-01-21        34.09
## 2020-01-22        34.02
## 2020-01-23        33.89
## [1] "-----"
## [1] "BAC"
##      BAC.Open BAC.High BAC.Low BAC.Close BAC.Volume BAC.Adjusted
```

```
## 2020-01-17    34.92    34.95    34.60    34.71    54156100    34.71
## 2020-01-21    34.42    34.52    34.22    34.26    50811000    34.26
## 2020-01-22    34.37    34.45    34.23    34.36    38879700    34.36
## 2020-01-23    34.09    34.26    33.73    34.12    45696200    34.12
## [1] "-----"
## [1] "RAD"
##          RAD.Open RAD.High RAD.Low RAD.Close RAD.Volume RAD.Adjusted
## 2020-01-17    12.96    13.48    12.45    12.60    4672500    12.60
## 2020-01-21    12.49    13.45    12.35    12.98    4854900    12.98
## 2020-01-22    12.97    13.28    12.72    12.86    2808500    12.86
## 2020-01-23    12.68    12.92    12.36    12.58    2453400    12.58
## [1] "-----"
## [1] "PFE"
##          PFE.Open PFE.High PFE.Low PFE.Close PFE.Volume PFE.Adjusted
## 2020-01-17    40.57    40.81    40.42    40.51    21901300    40.51
## 2020-01-21    40.38    40.66    40.19    40.34    21931400    40.34
## 2020-01-22    40.32    40.41    39.97    40.19    17170600    40.19
## 2020-01-23    40.19    40.83    40.13    40.71    25757600    40.71
## [1] "-----"
```

While all of our stocks have different starting dates, they all have the same end date which is the most recent closing date (it will change depending on when you run the code).

Not all data is relevant for our puposes. What we are really interested in are the adjusted prices which are given in the last column of each list object. The adjusted prices are adjusted for stock splits and dividends. The idea is that the adjusted prices represent the true return on an investment in a given stock.

Therefore, the next step is to get all the adjusted prices into a single object:

```
# Get the adjusted prices into a single object
P.adj <- lapply(P.list, function(p) p[,6])

# Merge the elements of the list
P <- Reduce(merge,P.adj)
names(P) <- tics

head(P,10)
```

```
##          AAPL      GE TWTR      BAC RAD      PFE
## 1980-01-02    NA 0.004071    NA 0.130777    NA 0.000681
## 1980-01-03    NA 0.004124    NA 0.129566    NA 0.000681
## 1980-01-04    NA 0.004259    NA 0.129566    NA 0.000706
## 1980-01-07    NA 0.004405    NA 0.130777    NA 0.000706
## 1980-01-08    NA 0.004562    NA 0.129566    NA 0.000740
## 1980-01-09    NA 0.004510    NA 0.135620    NA 0.000724
## 1980-01-10    NA 0.004531    NA 0.133198    NA 0.000715
## 1980-01-11    NA 0.004520    NA 0.133198    NA 0.000724
## 1980-01-14    NA 0.004520    NA 0.133198    NA 0.000724
## 1980-01-15    NA 0.004416    NA 0.134409    NA 0.000706
```

```
tail(P,10)
```

```
##          AAPL      GE TWTR      BAC RAD      PFE
## 2020-01-09 309.63 11.91 33.22 35.03 12.69 38.89
## 2020-01-10 310.33 11.67 32.78 34.74 12.04 39.49
## 2020-01-13 316.96 12.12 32.69 35.06 11.66 39.41
## 2020-01-14 312.68 12.03 32.82 35.32 12.28 40.07
```

```
## 2020-01-15 311.34 11.87 33.23 34.67 12.83 40.67
## 2020-01-16 315.24 11.84 34.19 34.72 13.00 40.61
## 2020-01-17 318.73 11.81 34.22 34.71 12.60 40.51
## 2020-01-21 316.57 11.66 34.09 34.26 12.98 40.34
## 2020-01-22 317.70 11.37 34.02 34.36 12.86 40.19
## 2020-01-23 319.23 11.77 33.89 34.12 12.58 40.71
```

Again, the P.adj “lapply” function could be written equivalently in a loop as follows:

```
# Get the adjusted prices into a single object
P.adj2 <- list(NA)
for (i in 1:n){
  P.adj2[[i]] <- P.list[[i]][,6]
}

P2 <- Reduce(merge,P.adj2)
names(P2) <- tics

head(P2,10)
```

```
##          AAPL      GE TWTR      BAC RAD      PFE
## 1980-01-02    NA 0.004071    NA 0.130777    NA 0.000681
## 1980-01-03    NA 0.004124    NA 0.129566    NA 0.000681
## 1980-01-04    NA 0.004259    NA 0.129566    NA 0.000706
## 1980-01-07    NA 0.004405    NA 0.130777    NA 0.000706
## 1980-01-08    NA 0.004562    NA 0.129566    NA 0.000740
## 1980-01-09    NA 0.004510    NA 0.135620    NA 0.000724
## 1980-01-10    NA 0.004531    NA 0.133198    NA 0.000715
## 1980-01-11    NA 0.004520    NA 0.133198    NA 0.000724
## 1980-01-14    NA 0.004520    NA 0.133198    NA 0.000724
## 1980-01-15    NA 0.004416    NA 0.134409    NA 0.000706
```

```
tail(P2,10)
```


```
##          AAPL      GE TWTR      BAC RAD      PFE
## 2020-01-09 309.63 11.91 33.22 35.03 12.69 38.89
## 2020-01-10 310.33 11.67 32.78 34.74 12.04 39.49
## 2020-01-13 316.96 12.12 32.69 35.06 11.66 39.41
## 2020-01-14 312.68 12.03 32.82 35.32 12.28 40.07
## 2020-01-15 311.34 11.87 33.23 34.67 12.83 40.67
## 2020-01-16 315.24 11.84 34.19 34.72 13.00 40.61
## 2020-01-17 318.73 11.81 34.22 34.71 12.60 40.51
## 2020-01-21 316.57 11.66 34.09 34.26 12.98 40.34
## 2020-01-22 317.70 11.37 34.02 34.36 12.86 40.19
## 2020-01-23 319.23 11.77 33.89 34.12 12.58 40.71
```

Now we have a data set of equity prices which contains time series observations in along the rows and the stocks across the columns. In the next section, we will prepare the data to form a set of efficient portfolios.

### 3 Macro Data

This section gives an example of how to get some interest rate or other macroeconomic data. The example we will use is the FRED database which is maintained by the Federal Reserve Bank of St. Louis and is available [here](#). Simply search for the data you want and the data code will be returned with the search.

```
#-----
#      PART 3: Getting interest rate data:
indices <- c("USD1MTD156N","DGS3MO")
ind.list <- lapply(indices, function(tic) get(getSymbols(tic, from = "1980-01-01", src = "FRED"))) )
IND <- Reduce(merge,ind.list)
plot(IND)
```



```
write.csv(data.frame(IND),file = "IND.csv")
```

## 4 Time-Series Data in R

This section provides some useful commands for handling time-series data in R. Note that there are other ways to handle time-series data.

Let's start by importing daily S&P 500 returns using the `quantmod` package. We'll start from January 1, 1990. We will use the 'SPDR S&P 500 ETF Trust ETF' to proxy the actual index.

```
## ----warning=FALSE,message=FALSE-----
library(quantmod)
library(lubridate)

##
## Attaching package: 'lubridate'

## The following object is masked from 'package:base':
##
##      date
```

```
library(plyr)
```

```
##  
## Attaching package: 'plyr'  
## The following object is masked from 'package:lubridate':  
##  
##     here
```

```
library(glmnet)
```

```
## Loading required package: Matrix  
## Loading required package: foreach  
## Loaded glmnet 2.0-18
```

```
library(PerformanceAnalytics)
```

```
##  
## Attaching package: 'PerformanceAnalytics'  
## The following object is masked from 'package:graphics':  
##  
##     legend
```

```
library(ggplot2)  
library(plotly)
```

```
##  
## Attaching package: 'plotly'  
## The following object is masked from 'package:ggplot2':  
##  
##     last_plot  
## The following objects are masked from 'package:plyr':  
##  
##     arrange, mutate, rename, summarise  
## The following object is masked from 'package:stats':  
##  
##     filter  
## The following object is masked from 'package:graphics':  
##  
##     layout
```

```
library(parallel)
```

```
rm(list = ls())
```

```
t1 <- "1990-01-01"  
v <- c("SPY")  
P.list <- lapply(v, function(sym) get(getSymbols(sym, from = t1)) )
```

Summarize the data.

```
head(P.list[[1]],10)
```

```
##           SPY.Open SPY.High  SPY.Low SPY.Close SPY.Volume SPY.Adjusted
```



```
## 1993-01-29 43.96875 43.96875 43.75000 43.93750 1003200 26.45393
## 1993-02-01 43.96875 44.25000 43.96875 44.25000 480500 26.64206
## 1993-02-02 44.21875 44.37500 44.12500 44.34375 201300 26.69851
## 1993-02-03 44.40625 44.84375 44.37500 44.81250 529400 26.98074
## 1993-02-04 44.96875 45.09375 44.46875 45.00000 531500 27.09362
## 1993-02-05 44.96875 45.06250 44.71875 44.96875 492100 27.07482
## 1993-02-08 44.96875 45.12500 44.90625 44.96875 596100 27.07482
## 1993-02-09 44.81250 44.81250 44.56250 44.65625 122100 26.88667
## 1993-02-10 44.65625 44.75000 44.53125 44.71875 379600 26.92429
## 1993-02-11 44.78125 45.12500 44.78125 44.93750 19500 27.05599
```

```
## -----
sapply(P.list,dim)
```

```
##      [,1]
## [1,] 6795
## [2,]    6
```

```
## -----
lapply(P.list, function(x) first(date(x)) )
```

```
## [[1]]
## [1] "1993-01-29"
```

```
## -----
P.list5 <- lapply(P.list, function(x) x[,5])
P.list6 <- lapply(P.list, function(x) x[,6])

## -----
P5 <- na.omit(Reduce(function(...) merge(...),P.list5 ))
P6 <- na.omit(Reduce(function(...) merge(...),P.list6 ))

head(P5, 5)
```

```
##          SPY.Volume
## 1993-01-29    1003200
## 1993-02-01     480500
## 1993-02-02     201300
## 1993-02-03     529400
## 1993-02-04     531500
```

```
head(P6, 5)
```

```
##          SPY.Adjusted
## 1993-01-29     26.45393
## 1993-02-01     26.64206
## 1993-02-02     26.69851
## 1993-02-03     26.98074
## 1993-02-04     27.09362
```

*P.5* and *P.6* are xts zoo objects containing the daily volume and adjusted price of *SPY*.

```
# adjust names
names(P5) <- names(P6) <- c("SPY")
names(P5) <- paste(names(P5),"vol",sep = "_")
head(P5,5)
```

```
##          SPY_vol
## 1993-01-29 1003200
```

```
## 1993-02-01 480500
## 1993-02-02 201300
## 1993-02-03 529400
## 1993-02-04 531500
```

```
head(P6,5)
```

```
##          SPY
## 1993-01-29 26.45393
## 1993-02-01 26.64206
## 1993-02-02 26.69851
## 1993-02-03 26.98074
## 1993-02-04 27.09362
```

Next, let's compute the returns to forecast.

```
## -----
R6 <- Return.calculate(P6)
head(R6, 5)
```

```
##          SPY
## 1993-01-29    NA
## 1993-02-01 0.007111495
## 1993-02-02 0.002118830
## 1993-02-03 0.010571190
## 1993-02-04 0.004183799
```

Let's say we wanted a rolling average of the returns to implement a momentum strategy. We can use the `rollapply()` function.

```
# add 5-day rolling average:
n.days <- 5
R6_roll <- rollapply(R6,n.days,mean)
names(R6_roll) <- paste(names(R6_roll),"_roll_",n.days,sep="")

R <- na.omit(merge(R6,R6_roll,P5))

SPY_lag <- stats::lag(R$SPY,1)
names(SPY_lag) <- "SPY_lag"
R <- na.omit(merge(SPY_lag,R))

## -----
range(date(R))
```

```
## [1] "1993-02-08" "2020-01-23"
```

```
## -----
cor(R)[,"SPY"]
```

```
##      SPY_lag      SPY  SPY_roll_5      SPY_vol
## -0.06156836  1.00000000  0.42700574 -0.07683575
```

```
## -----
cor(R)[,"SPY_lag"]
```

```
##      SPY_lag      SPY  SPY_roll_5      SPY_vol
##  1.00000000 -0.06156836  0.40392257 -0.10460493
```

```
# stack into a dataset rather than an xts object
ds <- data.frame(date = date(R),R)
```

```
rownames(ds) <- NULL
library("kableExtra")
```

```
## Warning: package 'kableExtra' was built under R version 3.6.2
```

```
kable(ds[c(1:20),], digits=4)
```

date	SPY_lag	SPY	SPY_rol_5	SPY_vol
1993-02-08	-0.0007	0.0000	0.0032	596100
1993-02-09	0.0000	-0.0069	0.0014	122100
1993-02-10	-0.0069	0.0014	-0.0004	379600
1993-02-11	0.0014	0.0049	-0.0003	19500
1993-02-12	0.0049	-0.0076	-0.0017	42500
1993-02-16	-0.0076	-0.0252	-0.0067	374800
1993-02-17	-0.0252	-0.0007	-0.0055	210900
1993-02-18	-0.0007	-0.0007	-0.0059	378100
1993-02-19	-0.0007	0.0036	-0.0061	34900
1993-02-22	0.0036	0.0036	-0.0039	513600
1993-02-23	0.0036	-0.0007	0.0010	373700
1993-02-24	-0.0007	0.0129	0.0037	26300
1993-02-25	0.0129	0.0021	0.0043	44500
1993-02-26	0.0021	0.0014	0.0039	66200
1993-03-01	0.0014	-0.0028	0.0026	66500
1993-03-02	-0.0028	0.0148	0.0057	182400
1993-03-03	0.0148	0.0042	0.0039	280100
1993-03-04	0.0042	-0.0055	0.0024	89500
1993-03-05	-0.0055	-0.0028	0.0016	40000
1993-03-08	-0.0028	0.0223	0.0066	50800


Now we have a dataset that we could use to predict returns, for example. If we wanted to extract a vector of the weeks of the data where each date corresponds to a Sunday, we could do the following. The same could be done for months.

```
## ----warning=FALSE-----
weeks <- date(unique(floor_date(ds$date, "week")))
weeks <- c(weeks, last(weeks) + weeks(1))

months <- date(unique(floor_date(ds$date, "month")))
months <- c(months, last(months) + months(1))
```

As a simple example, we could run a linear regression of the returns on the lagged returns. First, visualize a scatterplot. Note that we don't expect to see much here since returns are hard to predict!

```
scatter.smooth(x=ds$SPY_lag, y=ds$SPY, main="R_t ~ R_{t-1}") # scatterplot
```



Now run the model and output a summary of the results.

```
sp.mod <- lm(SPY ~ SPY_lag, data=ds)
print(sp.mod)
```

```
##
## Call:
## lm(formula = SPY ~ SPY_lag, data = ds)
##
## Coefficients:
## (Intercept)      SPY_lag
##    0.000461    -0.061568
```

```
summary(sp.mod)
```

```
##
## Call:
## lm(formula = SPY ~ SPY_lag, data = ds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.099820 -0.004687  0.000318  0.005284  0.143243
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.0004610  0.0001384   3.331 0.000869 ***
## SPY_lag      -0.0615683  0.0121154  -5.082 3.84e-07 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.01139 on 6787 degrees of freedom
## Multiple R-squared:  0.003791,    Adjusted R-squared:  0.003644
## F-statistic: 25.83 on 1 and 6787 DF,  p-value: 3.838e-07
```

While the lagged returns are statistically significant, they explain virtually none of the error as the  $adj-R^2 = 0.0036$  is tiny.

## 5 References